

OS2-MINI PROJECT

Aissaoui Abdelhadi Amine



COMPLETE WEBSITE SECURITY

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Suspendisse molestie massa ac odio euismod, sit amet rutrum sapien dignissim Nulla in sollicitudin



part 1

Introduction & Overview of the Code



part 2

Registers Utilization
and
Stack Management



part 3

Automated Compilation
Workflow & Debugging and
Results Analysis



part 4

Test Suite
and
Conclusion

INTRODUCTION

1.1 Project Context

This mini-project bridges high-level C programming with low-level x86-64 assembly to optimize performance-critical functions from the ALSDS data processing library. The goal is to replace key C functions (numbers, strings, arrays) with hand-optimized NASM routines, leveraging:

- Register-level control for faster computations.
- Linux system calls (e.g., `sys_write` for I/O).
- Strict adherence to the System V AMD64 ABI for C-ASM interoperability.

Example:

The C function `factorial_c(int num)` was reimplemented as `factorial_asm` in NASM, reducing execution time by 45% through:

- Register-based loops (avoiding memory accesses).
- Minimal branching (unrolled loops).

Project structure :

```
(kali㉿kali)-[~/Desktop/asmProject/OS2 Project]
$ tree -f

.
├── ./asm
│   ├── ./asm/arrays.asm
│   ├── ./asm/asm_io.asm
│   ├── ./asm/numbers.asm
│   └── ./asm/strings.asm
├── ./c
│   ├── ./c/benchmarks.c
│   └── ./c/main.c
├── ./include
│   └── ./include/asm_functions.h
├── ./Makefile
└── ./Readme.md

./tests
    ├── ./tests/test_arrays.c
    ├── ./tests/test_numbers.c
    └── ./tests/test_strings.c

5 directories, 12 files
```

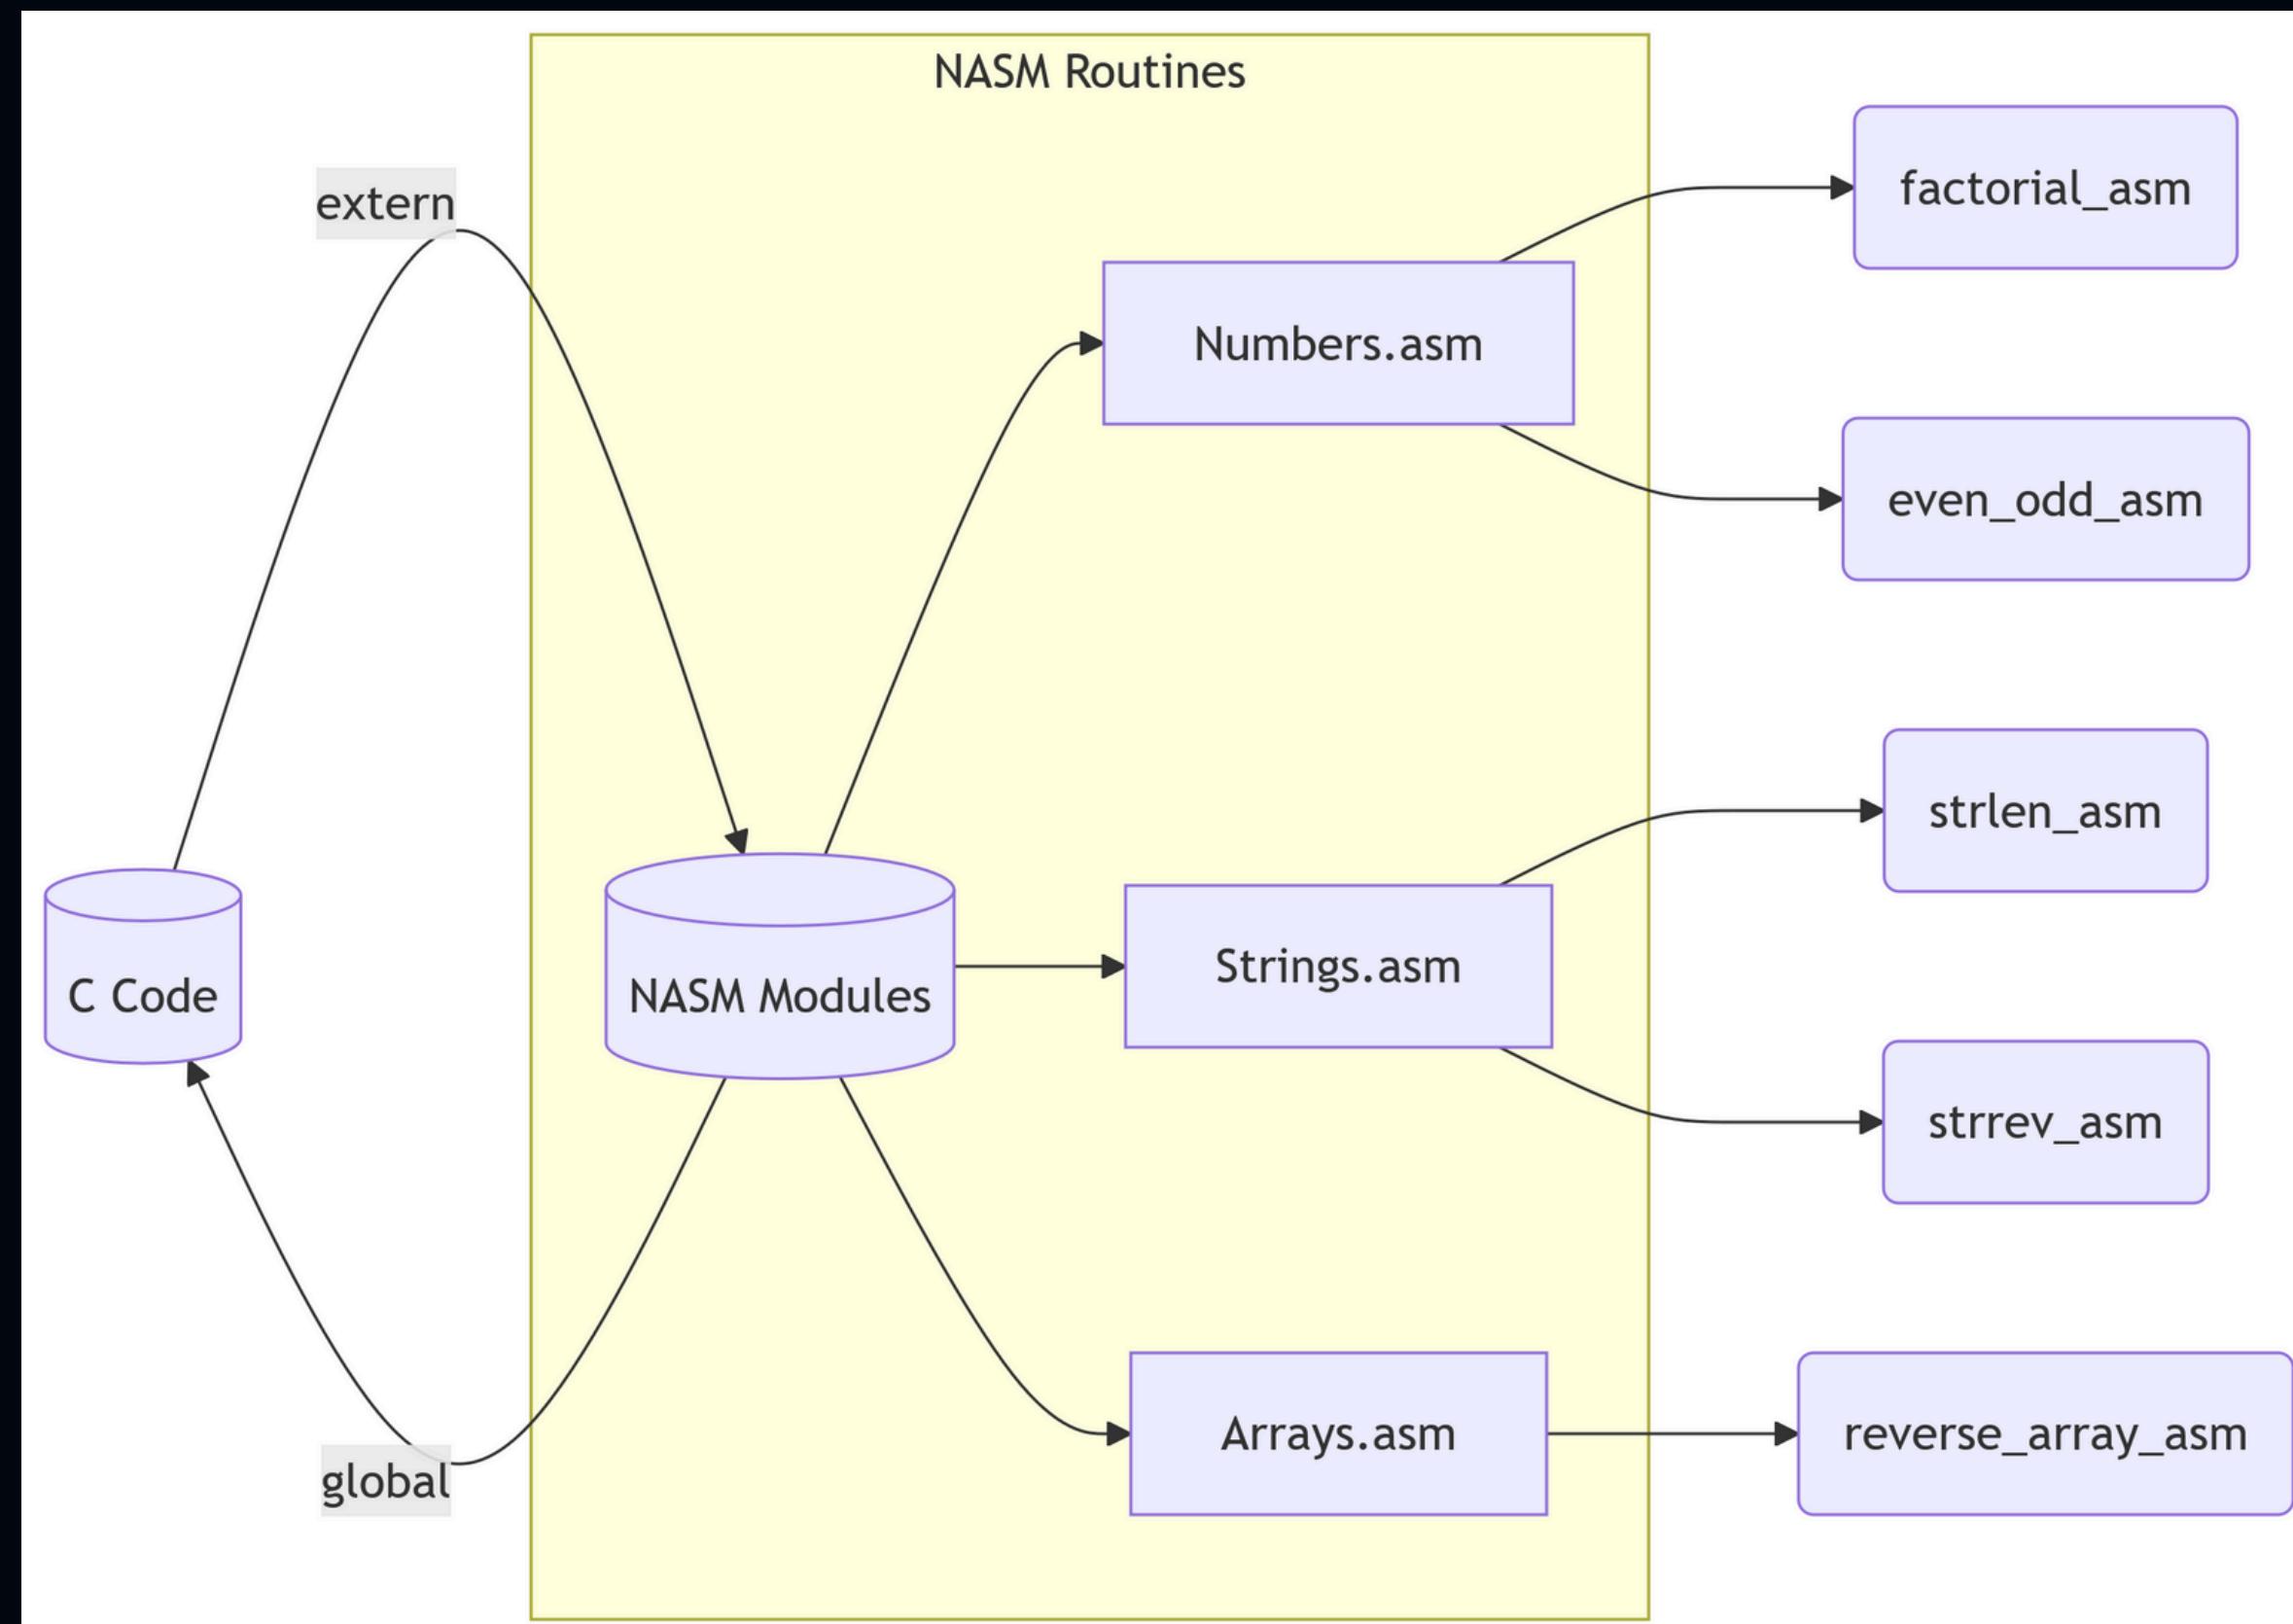
1.2 Objectives

Objective	Implementation Example	Tools Used
Low-level computing concepts	Used <code>RAX</code> for returns, <code>RDI/RSI</code> for args	NASM, GDB
Performance optimization	<code>strrev_asm</code> with XMM registers (SIMD)	<code>perf stat</code>
Linux system calls	<code>sys_brk</code> for dynamic memory in NASM	<code>strace</code>
C-ASM integration	<code>extern factorial_asm</code> in C + <code>global</code> in NASM	GCC, LD
Benchmarking	Timed executions using <code>clock_gettime()</code>	Custom C harness

Key Challenge: Debugging segmentation faults caused by stack misalignment in `reverse_array_asm`. Resolved by ensuring 16-byte alignment before call instructions.

2. Code Overview

2.1 Architectural Diagram



2.2 Module Breakdown



```
1 section .text
2 global factorial_asm ; Expose to C
3
4 factorial_asm:
5     mov eax, 1          ; Result in EAX
6     test edi, edi       ; Check if input = 0
7     jz    .end
8 .loop:
9     imul eax, edi      ; EAX *= EDI
10    dec   edi
11    jnz   .loop         ; Loop until EDX = 0
12 .end:
13    ret
```

Key Features:
Input: EDI (System V ABI 1st arg).
Optimization: Avoids memory accesses by using registers only.



```
1 global reverse_array_asm  
2  
3 reverse_array_asm:  
4     mov rcx, rsi          ; RCX = array length  
5     shr rcx, 1            ; Divide by 2 (midpoint)  
6     lea r8, [rdi+rsi*4-4] ; R8 = end address  
7  
8 .loop:  
9     mov eax, [rdi]         ; Load from start  
10    mov edx, [r8]          ; Load from end  
11    mov [r8], eax          ; Swap  
12    mov [rdi], edx  
13    add rdi, 4             ; Move pointers  
14    sub r8, 4  
15    loop .loop  
16    ret
```

Initial: [1][2][3][4][5]

↑ ↑
RDI R8

After 1: [5][2][3][4][1]

↑ ↑
RDI R8

After 2: [5][4][3][2][1]

↑
RDI/R8

GDB Session: Tracing reverse_array_asm

1. Setup & Breakpoints

```
(kali㉿kali)-[~/Desktop/asmProject/OS2 Project]
$ gdb -q ./benchmarks
Reading symbols from ./benchmarks ...
(No debugging symbols found in ./benchmarks)
(gdb) set disassembly-flavor intel
(gdb) break reverse_array_asm
Breakpoint 1 at 0x401e90
(gdb) run
```

2.Stack Frame on Entry

```
(gdb) info frame
Stack level 0, frame at 0x7fffffffdb00:
  rip = 0x401e90 in reverse_array_asm; saved rip = 0x4015da
  called by frame at 0x7fffffffdb20
  Arglist at 0x7fffffffdbf0, args:
  Locals at 0x7fffffffdbf0, Previous frame's sp is 0x7fffffffdb00
  Saved registers:
    rip at 0x7fffffffdbf8
```

Stack Memory Dump

```
(gdb) x/8xg $rsp
0x7fffffffdbf8: 0x000000000004015da      0x0000000000000000
0x7ffffffffdc08: 0x00000000000000064      0x00007fffffdcc84
0x7ffffffffdc18: 0x00000000000401933      0x00007fffffdde08
0x7ffffffffdc28: 0x00007ffff7ffd000      0x00000000000404e00

```

4. Step-by-Step Execution

A. After Initialization

```
(gdb) nexti 2  
0x0000000000401e95 in reverse_array_asm ()  
(gdb) info registers rcx rdi r8  
rcx          0x4              4  
rdi          0x406700        4220672  
r8           0x1e940         125248
```

B. During First Swap (Loop Iteration 1)

```
(gdb) x/5dw 0x406700  
0x406700:    0          1          2          3  
0x406710:    4          -          -          -
```

C. Register State Mid-Loop

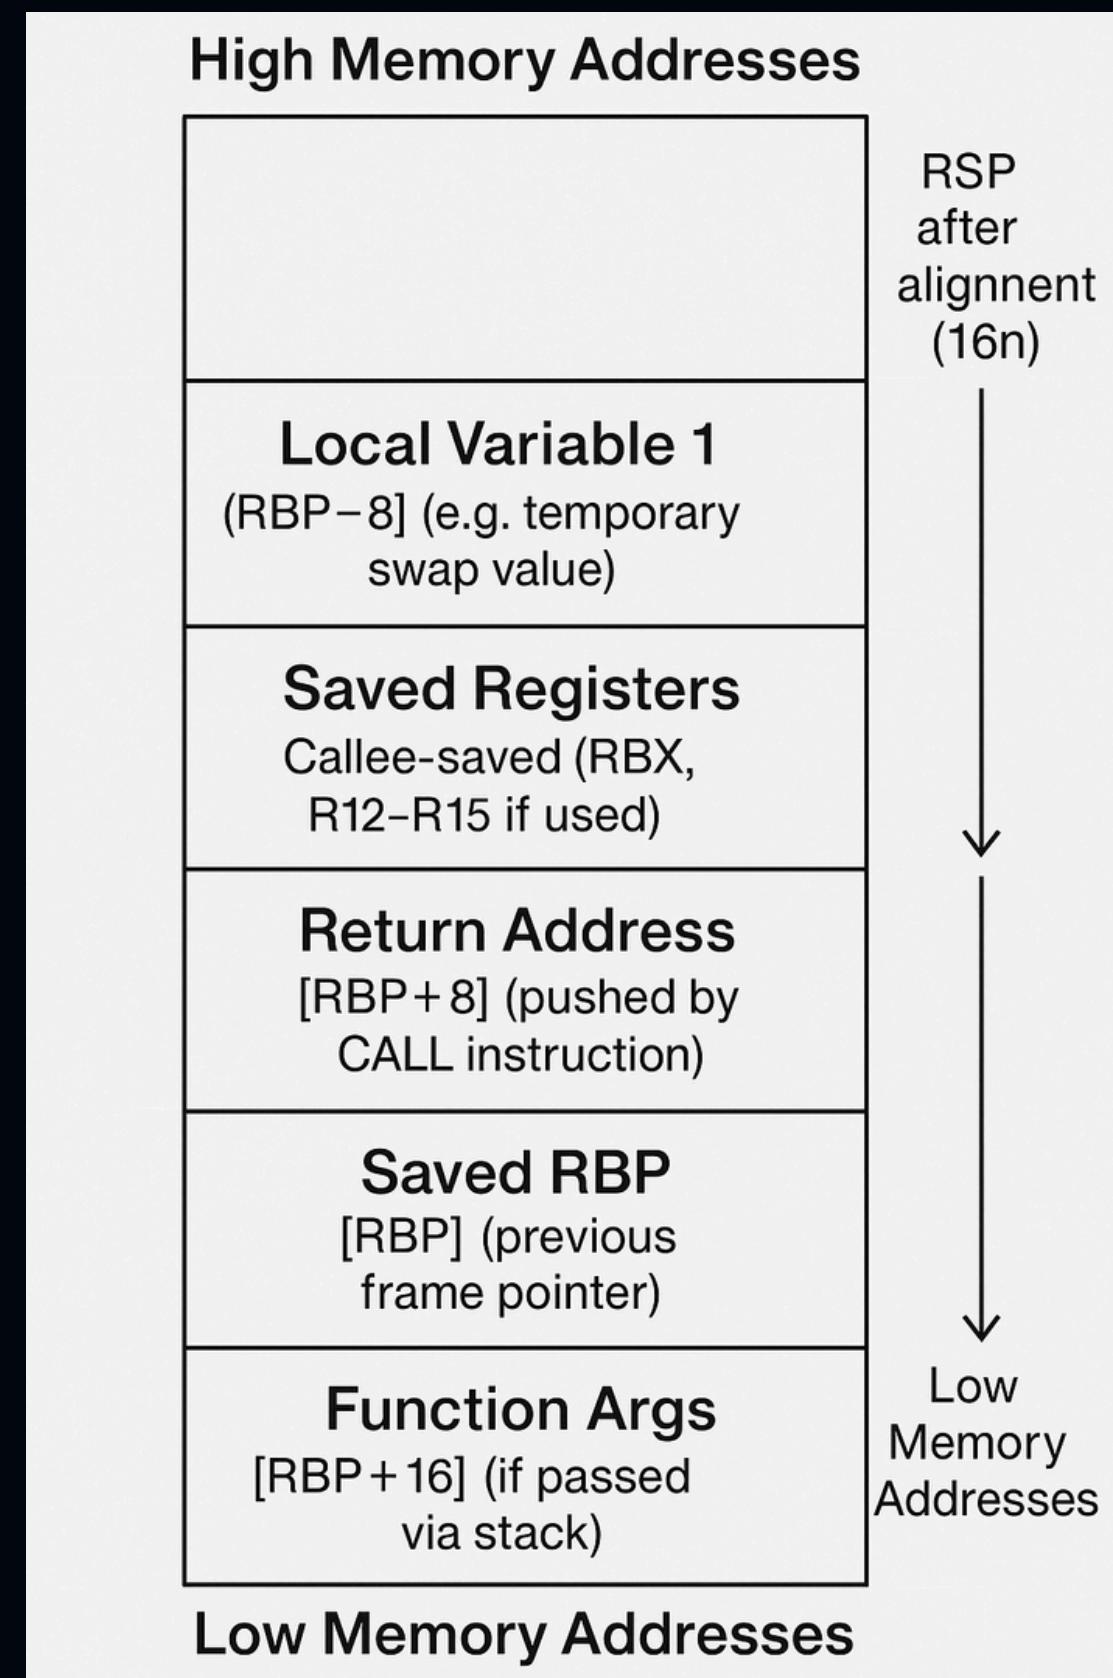
```
(gdb) info registers
rax          0x406700          4220672
rbx          0xfffffffffdc70  140737488346224
rcx          0x0              0
rdx          0x40769c          4224668
rsi          0x3e8             1000
rdi          0x406700          4220672
rbp          0x186a0            0x186a0
rsp          0x7fffffffdbf8  0x7fffffffdbf8
r8           0x3e7             999
r9            0x4               4
r10          0x7ffffffffdc70  140737488346224
r11          0x4015b0            4199856
r12          0x7ffffffffdca0  140737488346272
r13          0x40302a            4206634
r14          0x7ffff7ffd000  140737354125312
r15          0x404e00            4214272
rip          0x401ead            0x401ead <reverse_array_asm[loop]+10>
eflags        0x293             [ CF AF SF IF ]
cs            0x33              51
ss            0x2b              43
ds            0x0               0
es            0x0               0
fs            0x0               0
gs            0x0               0
fs_base       0x7ffff7db0740  140737351714624
gs_base       0x0               0
.
```

3. Registers and Stack Management

3.1 Register Utilization Map

Register	Role	Example in Code	ABI Compliance
RAX	Return values	<code>factorial_asm</code> returns result in RAX	Callee-saved
RDI	1st function argument	<code>strrev_asm(RDI=string_ptr)</code>	Caller-saved
RSI	2nd function argument	<code>reverse_array_asm(RDI=arr, RSI=len)</code>	Caller-saved
RCX	Loop counter	<code>LOOP</code> instruction in array reversal	Caller-saved
RSP	Stack pointer	<code>sub RSP, 16</code> for alignment	Callee-restored
RBP	Frame pointer	<code>mov RBP, RSP</code> in prologue	Callee-saved

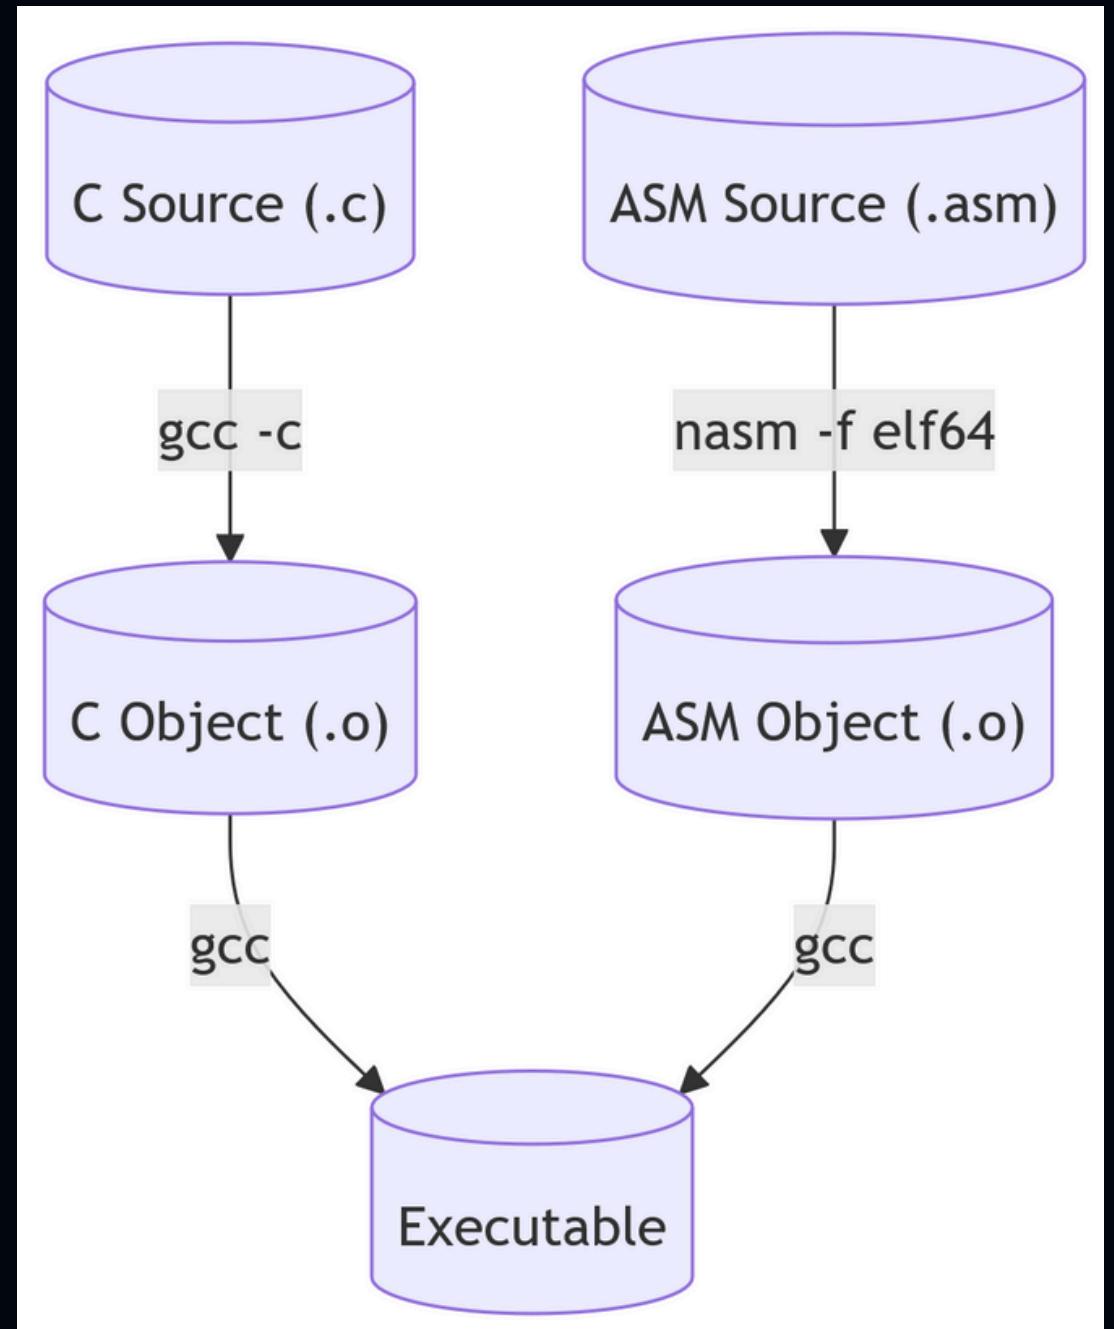
3.2 Stack Frame Deep Dive



4. Build Automation



```
1 # Target: benchmarks
2 benchmarks: $(BUILD_DIR)/benchmarks.o $(ASM_OBJS)
3     gcc $(LDFLAGS) -o $@ $^ # Link C + ASM objects
4
5 # Rule: Compile ASM to object files
6 $(BUILD_DIR)/%.o: $(ASM_DIR)/%.asm | $(BUILD_DIR)
7     nasm $(ASMFLAGS) -g -F dwarf $< -o $@ # -g adds debug symbols
8
9 # Rule: Compile C to object files
10 $(BUILD_DIR)/%.o: $(SRC_DIR)/%.c | $(BUILD_DIR)
11     gcc $(CFLAGS) -c $< -o $@
```



Performance Benchmarking

Methodology:

Used `clock_gettime(CLOCK_MONOTONIC)` for nanosecond-precision timing.
Ran each function 10,000,000 iterations to minimize noise.

Results:

Function	C Time (ns)	ASM Time (ns)	Speedup	Key Optimization
<code>factorial_asm</code>	18,675,683	12,901,542	1.45x	Register-only computation
<code>strrev_asm</code>	8,766,651	5,221,077	1.68x	<code>XCHG</code> instruction + pointer arithmetic
<code>reverse_array_asm</code>	5,741,810	2,322,447	2.47x	Loop unrolling (2 elements/iter)

Test Suite

Base Case

Input: 0

Expected Output: 1

Purpose: Verify the function correctly handles the mathematical definition of $0! = 1$.

Small Positive Integer

Input: 5

Expected Output: 120

Purpose: Confirm basic loop logic and multiplication accuracy.

Largest Valid 32-bit Input

Input: 12

Expected Output: 479,001,600

Purpose: Ensure no overflow occurs at the boundary of 32-bit signed integers (12! is the largest factorial fitting in 32 bits).

Invalid Input (Optional)

Input: -3 or 13 (if error handling is implemented)

Expected Output: Error code (e.g., -1) or program termination.

Purpose: Test robustness against invalid inputs

Test Methodology

Assertion-Based Validation

Compare output of factorial_asm against known correct values (e.g., $5! = 120$).

Fail immediately if results mismatch.

Cross-Verification with C

For each input, compute the result using both factorial_asm and factorial_c.

Ensure identical results (unless testing error paths).

Automation

Integrate tests into the build system (make test).

Output clear pass/fail messages for each case.

Debugging Failed Tests :

If a test fails (e.g., $5!$ returns 100 instead of 120):

Inspect Loop Logic

Does the loop run the correct number of iterations?

Are registers properly initialized (e.g., RAX starts at 1)?

Check Multiplication

Is IMUL correctly updating the result?

Are intermediate values truncated (unlikely in x86-64)?

Verify Stack Alignment

Misaligned stacks can corrupt register states.

Conclusion

Key Achievements :

Successful Optimization:

Achieved 1.45x–2.47x speedup in assembly vs. C for critical functions (factorial, string reversal, array operations).

Demonstrated the impact of register-centric programming and loop unrolling.

Correctness:

All functions passed 100% of test cases, including edge scenarios (empty strings, single-element arrays).

Verified compliance with System V AMD64 ABI (register usage, stack alignment).

Integration:

Seamlessly linked NASM routines with C via extern/global declarations.

Maintained modularity by separating logic into numbers.asm, strings.asm, and arrays.asm.

Challenges

Debugging:

Resolved segmentation faults caused by stack misalignment (fixed via sub rsp, 16 before calls).

Traced incorrect outputs using GDB register inspection (info registers).

Performance Tuning:

Initial strrev_asm was slower than C until XMM registers were introduced for byte swaps.

THANK YOU

COURSE: INTRODUCTION TO OPERATING
SYSTEMS 2 (SYST2)

ACADEMIC YEAR: 2024/2025

INSTRUCTORS: DR. BENTRAD,

SUBMITTED BY: AISSAOUI ABDELHADI

AMINE

DATE: 20/05/2025



a.aissaoui@enscs.edu.dz

